




*Ultra-low latency audio and
sensor processing
on the BeagleBone Black*




A project by
The Augmented Instruments Lab
at C4DM, Queen Mary University of London

<http://bela.io>

The Goal:

High-performance, self-contained
audio and sensor processing



The Goal: High-performance, self-contained audio and sensor processing



- Easy low-level hardware connectivity
- No OS = precise control of timing
- Very limited CPU (8-bit, 16MHz)
- Not good for audio processing
- Reasonable CPU (up to 1GHz ARM)
- High-level hardware (USB, network etc.)
- Limited low-level hardware
- Linux OS = high-latency / underruns
- Fast CPU
- High-level hardware (USB, network etc.)
- Arduino for low-level
- USB connection = high-latency, jitter
- Bulky, not self-contained



hardware



BeagleBone Black

1GHz ARM Cortex-A8

NEON vector floating point

PRU real-time microcontrollers

512MB RAM


Custom Bela Cape

Stereo audio in + out


Stereo 1.1W speaker amps

8x 16-bit analog in + out

16x digital in/out



features



1ms round-trip audio latency without underruns


High sensor bandwidth: digital I/Os sampled at 44.1kHz; analog I/Os sampled at 22.05kHz

Jitter-free alignment between audio and sensors



Hard real-time audio+sensor performance, but full Linux APIs still available

Programmable using **C/C++, Pd or Faust**

Designed for **musical instruments and live audio**




software



Xenomai Linux kernel

Debian distribution
Xenomai hard real-time
extensions




C++ programming API

Uses PRU for audio/sensors
Runs at higher priority
than kernel = *no dropouts*
Buffer sizes as small as **2**


Bela software



- Hard real-time environment using Xenomai Linux kernel extensions
- Use BeagleBone Programmable Realtime Unit (PRU) to write straight to hardware



- Sample all matrix ADCs and DACs at half audio rate (22.05kHz)
- Buffer sizes as small as 2 samples (90µs latency)



- **Speakers** with on-board amps
- **Audio In**
- **Audio Out**
- **16x digital I/O**
- **8x 16-bit analogue in (22.05kHz)**
- **8x 16-bit analogue out (22.05kHz)**

Find an interactive pin out diagram at <http://bela.io/belaDiagram>

Getting Started

bela.io/code/wiki

Materials

what you need to get started...

- **BeagleBone Black (BBB)**
- **Bela Cape**
- **SD card** with Bela image
- 3.5mm headphone jack **adapter cable**
- **Mini-USB cable** (to attach BBB to computer)
- Also useful for hardware hacking: **breadboard**, **jumper wires**, etc.

Step 1

install BBB drivers and Bela software

Install the BeagleBone Black drivers for your OS:


<http://bela.io/code/wiki> --> Getting Started

Bela code (for later):

<http://bela.io/code> --> Downloads --> bela-ableton-workshop.zip

Step 2: Access the IDE:

<http://192.168.7.2:3000>



The screenshot shows the Bela IDE interface. The main window has two tabs: "Bela IDE" and "BelaScope". The "Bela IDE" tab is active, displaying a code editor with C++ code. The code defines a class with setup and render methods. The "BelaScope" tab is visible in the background. On the right side, there is a vertical toolbar with icons for back, forward, search, settings, file, and help. Below the code editor, a terminal window shows the output of the build process:

```
Connected to the Bela IDE!
BBB date set to: Thu Feb 18 18:46:49 UTC 2016
Building OSCMessageFactory.cpp...
...done
Building OSCServer.cpp...
...done
```


API introduction

- In `render.cpp`....
- Three main functions:
- **setup()**
*runs once at the beginning, before audio starts
gives channel and sample rate info*
- **render()**
*called repeatedly by Bela system ("callback")
passes input and output buffers for audio and sensors*
- **cleanup()**
*runs once at end
release any resources you have used*
- bela.io/code/embedded Code docs


First test program

```
float gPhase; /* Phase of the oscillator (global variable) */  
  
void render(BeagleRTContext *context, void *userData)  
{  
    /* Iterate over the number of audio frames */  
    for(unsigned int n = 0; n < context->audioFrames; n++) {  
        /* Calculate the output sample based on the phase */  
        float out = 0.8f * sinf(gPhase);  
  
        /* Update the phase according to the frequency */  
        gPhase += 2.0 * M_PI * gFrequency * gInverseSampleRate;  
        if(gPhase > 2.0 * M_PI)  
            gPhase -= 2.0 * M_PI;  
  
        /* Store the output in every audio channel */  
        for(unsigned int channel = 0;  
            channel < context->audioChannels; channel++)  
            context->audioOut[n * context->audioChannels  
                            + channel] = out;  
    }  
}
```

One-dimensional array holding interleaved audio data



Access the IDE:
<http://192.168.7.2:3000>




Connect a Potentiometer

a.k.a. a “pot” or knob


Interactive pinout: <http://bela.io/belaDiagram>

The pot has 3 pins
5V and GND on the outside
Bela analog in in the middle



Connect an LED*

* Light-Emitting Diode



How to build other projects

1. **Web interface:** <http://192.168.7.2:3000>
Edit and compile code on the board
2. **Building scripts:**
 1. **Heavy Pd-to-C compiler** (<https://enzienaudio.com>)
Make audio patches in Pd-vanilla, translate to C and compile on board
 2. **Libpd**
Compile Pd patches without Heavy - access to more objects but not as fast, but good for prototyping
 3. **Faust**
Build online, export to C++, run on Bela

Bela and PureData

Heavy	libpd
Proprietary compiler, cloud-based, MIT non-commercial code	Free
Targets a variety of platforms (C, js, Unity, VST2)	Many ports (ofxpd, webpd ...)
Requires internet connection and local compiling (~1 minute)	Instantaneous (save the pd patch and restart)
Generates fast, optimized code, uses little CPU	It is just Pd (...)

libpd on Bela

How to run PureData patches on Bela with libpd :

1. Go to <http://bela.io/code/files> and download the bela-ableton-2016-04-12.zip archive
2. Unzip the archive into a convenient location and open a terminal window
3. Navigate into the scripts/ folder and run
`./run_pd_libpd.sh/projects/heavy/pd/demo-track/`
4. Type "yes" and you should hear something

Bela and Faust

- Today: you will have to download the C++ file generated by the <http://faust.grame.fr/onlinecompiler/> (after setting the -i flag), save it on your computer and target it with the build_project.sh script, as in:

```
/path/to/bela/repo/scripts/build_project.sh /path/  
to/faust/file/CppCode.cpp
```

```
freq = hslider("[1]Frequency[BELA:ANALOG_0]",  
440,460,1500,1):smooth(0.999);  
pressure = hslider("[2]Pressure[style:knob][BELA:ANALOG_4]", 0.96, 0.2,  
2.0, 0.01):smooth(0.999):min(0.99):max(0.2);  
gate = hslider("[0]ON/OFF (ASR Envelope)[BELA:DIGITAL_0]",0,0,1,1);
```


Bela Cape

I2C and GPIO

Audio In

Audio Out
(headphone)

Speakers



8-ch. 16-bit ADC


8-ch. 16-bit DAC

Connect a Potentiometer

a.k.a. a “pot” or knob


Interactive pinout: <http://bela.io/belaDiagram>

The pot has 3 pins
5V and GND on the outside
Bela analog in in the middle




Connect a LDR/FSR*

* Light-Dependent Resistor / Force-Sensing Resistor





Connect an LED*

* Light-Emitting Diode



Green pins can
be used for digital
I/O



API introduction


- In `render.cpp`....
- Three main functions:
- `setup()`
*runs once at the beginning, before audio starts
gives channel and sample rate info*
- `render()`
*called repeatedly by Bela system ("callback")
passes input and output buffers for audio and sensors*
- `cleanup()`
*runs once at end
release any resources you have used*

Real-time audio

- Suppose we have code that runs **offline**
 - ▶ (non-real time)
- Our goal is to re-implement it **online** (real time)
 - ▶ Generate audio **as we need it!**
 - ▶ Why couldn't we just generate it all in advance, and then play it when we need it?
- Digital audio is composed of **samples**
 - ▶ 44100 samples per second in our example
 - ▶ That means we need a new sample every $1/44100$ seconds (about every $23\mu\text{s}$)
 - ▶ So option #1 is to run a short bit of code every sample whenever we want to know what to play next
 - ▶ What might be some drawbacks of this approach?
 - Can we guarantee we'll be ready for each new sample?


Block-based processing

- Option #2: Process in **blocks** of several samples
 - ▶ Basic idea: generate enough samples to get through the next few milliseconds
 - ▶ Typical **block sizes**: 32 to 1024 samples
 - Usually a power of 2 for reasons having to do with hardware
 - ▶ While the audio hardware is busy playing one block, we can start calculating the next one so it's ready on time:



Block-based processing

- Option #2: Process in **blocks** of several samples
 - ▶ Basic idea: generate enough samples to get through the next few milliseconds
 - ▶ Typical **block sizes**: 32 to 1024 samples
 - Usually a power of 2 for reasons having to do with hardware
 - ▶ While the audio hardware is busy playing one block, we can start calculating the next one so it's ready on time:




Block-based processing


- Advantages of blocks over individual samples
 - ▶ We need to run our function less often
 - ▶ We always generate one block ahead of what is actually playing
 - ▶ Suppose one block of samples lasts 5ms, and running our code takes 1ms
 - ▶ Now, we can tolerate a delay of up to 4ms if the OS is busy with other tasks
 - ▶ Larger block size = can tolerate more variation in timing
- What is the disadvantage?
 - ▶ Latency (delay)

Latency


- Primary tradeoff for buffering: **latency**
 - ▶ There will be a **delay** from input to output
- Let's consider a full-duplex system (in and out)
 - ▶ Which are the sources of latency? **We have been writing this**



Latency: the role of buffering




- Block-based processing introduces latency
 - ▶ This is in addition to whatever was generated by $H(z)$
- On input side: how is a block of samples created?
 - ▶ For block of size N: we wait until N samples have arrived from ADC....




In other words: first sample in the block is already N samples old by the time we get it


Latency: the role of buffering



- On output side: how is a block played by DAC?
 - ▶ We can only **start** playing once the block arrives!
 - ▶ So how long until the last sample is played?
 - **N samples** after the the block is sent to the hardware Block



Buffering illustration



- At any given time, we are reading from ADC, processing, and writing, two buffers to the output
1. First we fill up a buffer of samples
- Total latency is 2x buffer length

API introduction


```
void render(BeagleRTContext *context, void *userData)
```

- Sensor ("matrix" = ADC+DAC) data is gathered **automatically** alongside audio
- Audio runs at **44.1kHz**; sensor data at **22.05kHz**
- **context** holds buffers plus information on number of frames and other info
- Your job as programmer: render one buffer of audio and sensors and finish as soon as possible!
- API documentation: <http://beaglert.cc>

First test program


```
float gPhase; /* Phase of the oscillator (global variable) */  
  
void render(BeagleRTContext *context, void *userData)  
{  
    /* Iterate over the number of audio frames */  
    for(unsigned int n = 0; n < context->audioFrames; n++) {  
        /* Calculate the output sample based on the phase */  
        float out = 0.8f * sinf(gPhase);  
  
        /* Update the phase according to the frequency */  
        gPhase += 2.0 * M_PI * gFrequency * gInverseSampleRate;  
        if(gPhase > 2.0 * M_PI)  
            gPhase -= 2.0 * M_PI;  
  
        /* Store the output in every audio channel */  
        for(unsigned int channel = 0;  
            channel < context->audioChannels; channel++)  
            context->audioOut[n * context->audioChannels  
                            + channel] = out;  
    }  
}
```

One-dimensional array holding interleaved audio data



Interleaving


- Two ways for multichannel audio to be stored
 - ▶ Way 1: Separate memory buffers per channel



- This is known as **non-interleaved** format
- Typically presented in C as a two-dimensional array:
`float **sampleBuffers`

- ▶ Way 2: **One memory buffer for all channels**

- Alternating data between channels



- This is known as **interleaved** format
- Typically presented in C as a one-dimensional array:
`float *sampleBuffer`

Interleaving

- We accessed non-interleaved data like this:

- ▶

```
float in = sampleBuffers[channel][n];
```

- How do we do the same thing with interleaving?

- ▶

```
float in = sampleBuffers[***what goes here?***];
```

- ▶ What else do we need to know?
 - Number of channels



- ▶

```
float in = sampleBuffers[numChannels*n + channel];
```


- ▶ Each sample advances **numChannels** in the buffer

- ▶ The **offset** tells us which channel we're reading

First test program

```
float gPhase; /* Phase of the oscillator (global variable) */  
  
void render(BeagleRTContext *context, void *userData)  
{  
    /* Iterate over the number of audio frames */  
    for(unsigned int n = 0; n < context->audioFrames; n++) {  
        /* Calculate the output sample based on the phase */  
        float out = 0.8f * sinf(gPhase);  
  
        /* Update the phase according to the frequency */  
        gPhase += 2.0 * M_PI * gFrequency * gInverseSampleRate;  
        if(gPhase > 2.0 * M_PI)  
            gPhase -= 2.0 * M_PI;  
  
        /* Store the output in every audio channel */  
        for(unsigned int channel = 0;  
            channel < context->audioChannels; channel++)  
            context->audioOut[n * context->audioChannels  
                            + channel] = out;  
    }  
}
```

One-dimensional array holding interleaved audio data




Blocks and phase: task


- Need to **preserve state** between calls to `render()`
 - ▶ When you call `render()` a second time, it should **remember where it left off** the first time
 - ▶ But local variables in the function all disappear when the function returns!
 - ▶ Solution: use **global variables** to save the state
 - Okay, cleaner solutions exist: keep a structure that you pass by pointer as an argument to `render()`. Save your state there.
 - Or in C++, use **instance variables** (variables that are declared in the class rather than within a method). But we'll save that for later.

Blocks and phase

- If we don't store phase in a global variable, we get:



- But what we want is this:



First test program

This remembers where we left off

```
float gPhase; /* Phase of the oscillator (global variable)
```


```
void render(BeagleRTContext *context, void *userData)
{
    /* Iterate over the number of audio frames */
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        /* Calculate the output sample based on the phase */
        float out = 0.8f * sinf(gPhase);

        /* Update the phase according to the frequency */
        gPhase += 2.0 * M_PI * gFrequency * gInverseSampleRate;
        if(gPhase > 2.0 * M_PI)
            gPhase -= 2.0 * M_PI;

        /* Store the output in every audio channel */
        for(unsigned int channel = 0;
            channel < context->audioChannels; channel++)
            context->audioOut[n * context->audioChannels
                + channel] = out;
    }
}
```

This updates the phase each sample and keeps it in the 0 to 2π range

Analog input data format



- Data type is `float`: just like audio
 - But range is 0.0 to 1.0
 - This is internally converted from raw values 0 to 65535
 - Compare this to audio, which is -1.0 to 1.0

Analog input

```
float gPhase;
float gInverseSampleRate;           /* Pre-calculated for convenience */
int gAudioFramesPerAnalogFrame;

extern int gSensorInputFrequency;   /* Which analog pin controls frequency */
extern int gSensorInputAmplitude;  /* Which analog pin controls amplitude */

void render(BeagleRTContext *context, void *userData)
{
    float frequency = 440.0;
    float amplitude = 0.8;

    for(unsigned int n = 0; n < context->audioFrames; n++) {
        /* There are twice as many audio frames as matrix frames since
         * audio sample rate is twice as high */
        if(!(n % gAudioFramesPerAnalogFrame)) {
            /* Every other audio sample: update frequency and amplitude */
            frequency = map(analogReadFrame(context,
                                              n/gAudioFramesPerAnalogFrame,
                                              gSensorInputFrequency),
                            0, 1, 100, 1000);
            amplitude = analogReadFrame(context,
                                         n/gAudioFramesPerAnalogFrame,
                                         gSensorInputAmplitude);
        }

        float out = amplitude * sinf(gPhase);

        for(unsigned int channel = 0; channel < context->audioChannels; channel++)
            context->audioOut[n * context->audioChannels + channel] = out;

        gPhase += 2.0 * M_PI * frequency * gInverseSampleRate;
        if(gPhase > 2.0 * M_PI)
            gPhase -= 2.0 * M_PI;
    }
}
```

This runs **every other sample**

Read the **analog input** at the specified **frame**


Map the 0-1 input range to a frequency range

Digital I/O

```
void render(BeagleRTContext *context, void *userData)
{
    static int count = 0; // counts elapsed samples
    float interval = 0.5; // how often to toggle the LED (in seconds)
    static int status = GPIO_LOW;

    for(unsigned int n = 0; n < context->digitalFrames; n++) {
        /* Check if enough samples have elapsed that it's time to
         * blink to the LED */
        if(count == context->digitalSampleRate * interval) {
            count = 0; // reset the counter
            if(status == GPIO_LOW) {
                /* Toggle the LED */
                digitalWriteFrame(context, n, P8_07, status);
                status = GPIO_HIGH;
            }
            else {
                /* Toggle the LED */
                digitalWriteFrame(context, n, P8_07, status);
                status = GPIO_LOW;
            }
        }
        /* Increment the count once per frame */
        count++;
    }
}
```

To manage timing, **count samples** rather than using delays



This runs **once per digital frame**

Write the **digital output** at the specified **frame**



Stay tuned! Join the announcement list at
<http://bela.io>